

Lending Club Default Analysis

The analysis is divided into four main parts:

1. Data understanding
2. Data cleaning (cleaning missing values, removing redundant columns etc.)
3. Data Analysis
4. Recommendations

```
In [14]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

data = pd.read_csv(r'C:\Users\91808\Downloads\loan.csv')
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 39717 entries, 0 to 39716
Columns: 111 entries, id to total_il_high_credit_limit
dtypes: float64(74), int64(13), object(24)
memory usage: 33.6+ MB
```

C:\Users\91808\AppData\Local\Temp\ipykernel_2656\286541020.py:6: DtypeWarning: Columns (47) have mixed types. Specify dtype option on import or set low_memory=False.

```
data = pd.read_csv(r'C:\Users\91808\Downloads\loan.csv')
```

Data Understanding

```
In [15]: # Let's look at the first few rows of the df
data.head()
```

Out[15]:

	id	member_id	loan_amnt	funded_amnt	funded_amnt_inv	term	int_rate	installment	grade	sub_grade	...	num_tl_90g_dpd
0	1077501	1296599	5000	5000	4975.0	36 months	10.65%	162.87	B	B2	...	
1	1077430	1314167	2500	2500	2500.0	60 months	15.27%	59.83	C	C4	...	
2	1077175	1313524	2400	2400	2400.0	36 months	15.96%	84.33	C	C5	...	
3	1076863	1277178	10000	10000	10000.0	36 months	13.49%	339.31	C	C1	...	
4	1075358	1311748	3000	3000	3000.0	60 months	12.69%	67.79	B	B5	...	

5 rows × 111 columns

```
In [16]: # Looking at all the column names
data.columns
```

```
Out[16]: Index(['id', 'member_id', 'loan_amnt', 'funded_amnt', 'funded_amnt_inv',
               'term', 'int_rate', 'installment', 'grade', 'sub_grade',
               ...,
               'num_tl_90g_dpd_24m', 'num_tl_op_past_12m', 'pct_tl_nvr_dlq',
               'percent_bc_gt_75', 'pub_rec_bankruptcies', 'tax_liens',
               'tot_hi_cred_lim', 'total_bal_ex_mort', 'total_bc_limit',
               'total_il_high_credit_limit'],
              dtype='object', length=111)
```

Some of the important columns in the dataset are loan_amount, term, interest rate, grade, sub grade, annual income, purpose of the loan etc.

The **target variable**, which we want to compare across the independent variables, is loan status. The strategy is to figure out compare the average default rates across various independent variables and identify the ones that affect default rate the most.

Data Cleaning

Some columns have a large number of missing values, let's first fix the missing values and then check for other types of data quality problems.

```
In [17]: # summarising number of missing values in each column
data.isnull().sum()
```

```
Out[17]: id                0
member_id                0
loan_amnt                0
funded_amnt              0
funded_amnt_inv          0

...
tax_liens                39
tot_hi_cred_lim          39717
total_bal_ex_mort        39717
total_bc_limit           39717
total_il_high_credit_limit 39717
Length: 111, dtype: int64
```

```
In [22]: # percentage of missing values in each column
round(data.isnull().sum()/len(data.index), 2)*100
```

```
Out[22]: id                0.0
member_id                0.0
loan_amnt                0.0
funded_amnt              0.0
funded_amnt_inv          0.0

...
tax_liens                0.0
tot_hi_cred_lim          100.0
total_bal_ex_mort        100.0
total_bc_limit           100.0
total_il_high_credit_limit 100.0
Length: 111, dtype: float64
```

You can see that many columns have 100% missing values, some have 65%, 33% etc. First, let's get rid of the columns having 100% missing values.

```
In [510... # removing the columns having more than 90% missing values
missing_columns = loan.columns[100*(loan.isnull().sum()/len(loan.index)) > 90]
print(missing_columns)
```

```
Index(['mths_since_last_record', 'next_pymnt_d', 'mths_since_last_major_derog',
      'annual_inc_joint', 'dti_joint', 'verification_status_joint',
      'tot_coll_amt', 'tot_cur_bal', 'open_acc_6m', 'open_il_6m',
      'open_il_12m', 'open_il_24m', 'mths_since_rcnt_il', 'total_bal_il',
      'il_util', 'open_rv_12m', 'open_rv_24m', 'max_bal_bc', 'all_util',
      'total_rev_hi_lim', 'inq_fi', 'total_cu_tl', 'inq_last_12m',
      'acc_open_past_24mths', 'avg_cur_bal', 'bc_open_to_buy', 'bc_util',
      'mo_sin_old_il_acct', 'mo_sin_old_rev_tl_op', 'mo_sin_rcnt_rev_tl_op',
      'mo_sin_rcnt_tl', 'mort_acc', 'mths_since_recent_bc',
      'mths_since_recent_bc_dlq', 'mths_since_recent_inq',
      'mths_since_recent_revol_delinq', 'num_accts_ever_120_pd',
      'num_actv_bc_tl', 'num_actv_rev_tl', 'num_bc_sats', 'num_bc_tl',
      'num_il_tl', 'num_op_rev_tl', 'num_rev_accts', 'num_rev_tl_bal_gt_0',
      'num_sats', 'num_tl_120dpd_2m', 'num_tl_30dpd', 'num_tl_90g_dpd_24m',
      'num_tl_op_past_12m', 'pct_tl_nvr_dlq', 'percent_bc_gt_75',
      'tot_hi_cred_lim', 'total_bal_ex_mort', 'total_bc_limit',
      'total_il_high_credit_limit'],
      dtype='object')
```

```
In [511... loan = loan.drop(missing_columns, axis=1)
print(loan.shape)
```

```
(39717, 55)
```

```
In [512... # summarise number of missing values again
100*(loan.isnull().sum()/len(loan.index))
```

```
Out[512]: id                0.000000
          member_id         0.000000
          loan_amnt          0.000000
          funded_amnt        0.000000
          funded_amnt_inv    0.000000
          term                0.000000
          int_rate            0.000000
          installment        0.000000
          grade               0.000000
          sub_grade           0.000000
          emp_title           6.191303
          emp_length          2.706650
          home_ownership      0.000000
          annual_inc          0.000000
          verification_status  0.000000
          issue_d             0.000000
          loan_status         0.000000
          pymnt_plan          0.000000
          url                 0.000000
          desc                32.580507
          purpose             0.000000
          title               0.027696
          zip_code            0.000000
          addr_state          0.000000
          dti                 0.000000
          delinq_2yrs         0.000000
          earliest_cr_line    0.000000
          inq_last_6mths      0.000000
          mths_since_last_delinq 64.662487
          open_acc            0.000000
          pub_rec             0.000000
          revol_bal           0.000000
          revol_util          0.125891
          total_acc           0.000000
          initial_list_status  0.000000
          out_prncp           0.000000
          out_prncp_inv       0.000000
          total_pymnt         0.000000
          total_pymnt_inv     0.000000
          total_rec_prncp     0.000000
```

```

total_rec_int          0.000000
total_rec_late_fee     0.000000
recoveries             0.000000
collection_recovery_fee 0.000000
last_pymnt_d           0.178765
last_pymnt_amnt        0.000000
last_credit_pull_d     0.005036
collections_12_mths_ex_med 0.140998
policy_code            0.000000
application_type        0.000000
acc_now_delinq          0.000000
chargeoff_within_12_mths 0.140998
delinq_amnt            0.000000
pub_rec_bankruptcies   1.754916
tax_liens              0.098195
dtype: float64

```

```

In [513]: # There are now 2 columns having approx 32 and 64% missing values -
# description and months since last delinquent

# Let's have a look at a few entries in the columns
loan.loc[:, ['desc', 'mths_since_last_delinq']].head()

```

```

Out[513]:

```

	desc	mths_since_last_delinq
0	Borrower added on 12/22/11 > I need to upgra...	NaN
1	Borrower added on 12/22/11 > I plan to use t...	NaN
2	NaN	NaN
3	Borrower added on 12/21/11 > to pay for prop...	35.0
4	Borrower added on 12/21/11 > I plan on combi...	38.0

The column description contains the comments the applicant had written while applying for the loan. Although one can use some text analysis techniques to derive new features from this column (such as sentiment, number of positive/negative words etc.), we will not use this column in this analysis.

Secondly, months since last delinquent represents the number months passed since the person last fell into the 90 DPD group. There is an important reason we shouldn't use this column in analysis - since at the time of loan application, we will not have this data (it gets generated months after the loan has been approved), it cannot be used as a predictor of default at the time of loan approval.

Thus let's drop the two columns.

```
In [514... # dropping the two columns  
loan = loan.drop(['desc', 'mths_since_last_delinq'], axis=1)
```

```
In [515... # summarise number of missing values again  
100*(loan.isnull().sum()/len(loan.index))
```

```
Out[515]: id                0.000000
          member_id         0.000000
          loan_amnt         0.000000
          funded_amnt       0.000000
          funded_amnt_inv   0.000000
          term              0.000000
          int_rate          0.000000
          installment       0.000000
          grade            0.000000
          sub_grade         0.000000
          emp_title         6.191303
          emp_length        2.706650
          home_ownership    0.000000
          annual_inc        0.000000
          verification_status 0.000000
          issue_d           0.000000
          loan_status       0.000000
          pymnt_plan        0.000000
          url              0.000000
          purpose          0.000000
          title            0.027696
          zip_code         0.000000
          addr_state       0.000000
          dti              0.000000
          delinq_2yrs       0.000000
          earliest_cr_line  0.000000
          inq_last_6mths    0.000000
          open_acc          0.000000
          pub_rec           0.000000
          revol_bal        0.000000
          revol_util        0.125891
          total_acc         0.000000
          initial_list_status 0.000000
          out_prncp         0.000000
          out_prncp_inv     0.000000
          total_pymnt       0.000000
          total_pymnt_inv   0.000000
          total_rec_prncp   0.000000
          total_rec_int     0.000000
          total_rec_late_fee 0.000000
```



```
recoveries          0.000000
collection_recovery_fee  0.000000
last_pymnt_d        0.178765
last_pymnt_amnt     0.000000
last_credit_pull_d   0.005036
collections_12_mths_ex_med  0.140998
policy_code         0.000000
application_type     0.000000
acc_now_delinq       0.000000
chargeoff_within_12_mths  0.140998
delinq_amnt         0.000000
pub_rec_bankruptcies  1.754916
tax_liens           0.098195
dtype: float64
```

There are some more columns with missing values, but let's ignore them for now (since we are not doing any modeling, we don't need to impute all missing values anyway).

But let's check whether some rows have a large number of missing values.

```
In [516... # missing values in rows
loan.isnull().sum(axis=1)
```

```
Out[516]: 0      1
          1      0
          2      1
          3      0
          4      0
          5      0
          6      0
          7      0
          8      1
          9      0
         10      0
         11      0
         12      0
         13      0
         14      0
         15      0
         16      0
         17      0
         18      0
         19      0
         20      0
         21      0
         22      0
         23      0
         24      0
         25      0
         26      1
         27      0
         28      0
         29      0
          ..
        39687    4
        39688    4
        39689    4
        39690    4
        39691    4
        39692    4
        39693    4
        39694    4
        39695    4
```

```
39696    4
39697    4
39698    4
39699    4
39700    5
39701    4
39702    4
39703    4
39704    5
39705    4
39706    5
39707    4
39708    4
39709    4
39710    4
39711    4
39712    4
39713    4
39714    5
39715    5
39716    4
```

Length: 39717, dtype: int64

```
In [517... # checking whether some rows have more than 5 missing values
len([loan[loan.isnull().sum(axis=1) > 5].index])
```

Out[517]: 0

The data looks clean by and large. Let's also check whether all columns are in the correct format.

```
In [518... loan.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 39717 entries, 0 to 39716
Data columns (total 53 columns):
id                39717 non-null int64
member_id        39717 non-null int64
loan_amnt        39717 non-null int64
funded_amnt      39717 non-null int64
funded_amnt_inv  39717 non-null float64
term            39717 non-null object
int_rate        39717 non-null object
installment     39717 non-null float64
grade          39717 non-null object
sub_grade      39717 non-null object
emp_title      37258 non-null object
emp_length     38642 non-null object
home_ownership 39717 non-null object
annual_inc     39717 non-null float64
verification_status 39717 non-null object
issue_d        39717 non-null object
loan_status    39717 non-null object
pymnt_plan     39717 non-null object
url            39717 non-null object
purpose        39717 non-null object
title          39706 non-null object
zip_code       39717 non-null object
addr_state     39717 non-null object
dti            39717 non-null float64
delinq_2yrs    39717 non-null int64
earliest_cr_line 39717 non-null object
inq_last_6mths 39717 non-null int64
open_acc       39717 non-null int64
pub_rec        39717 non-null int64
revol_bal      39717 non-null int64
revol_util     39667 non-null object
total_acc      39717 non-null int64
initial_list_status 39717 non-null object
out_prncp      39717 non-null float64
out_prncp_inv  39717 non-null float64
total_pymnt    39717 non-null float64
total_pymnt_inv 39717 non-null float64

```

```

total_rec_prncp      39717 non-null float64
total_rec_int        39717 non-null float64
total_rec_late_fee   39717 non-null float64
recoveries           39717 non-null float64
collection_recovery_fee 39717 non-null float64
last_pymnt_d         39646 non-null object
last_pymnt_amnt      39717 non-null float64
last_credit_pull_d   39715 non-null object
collections_12_mths_ex_med 39661 non-null float64
policy_code          39717 non-null int64
application_type     39717 non-null object
acc_now_delinq       39717 non-null int64
chargeoff_within_12_mths 39661 non-null float64
delinq_amnt          39717 non-null int64
pub_rec_bankruptcies 39020 non-null float64
tax_liens            39678 non-null float64
dtypes: float64(18), int64(13), object(22)
memory usage: 16.1+ MB

```

```

In [519... # The column int_rate is character type, let's convert it to float
loan['int_rate'] = loan['int_rate'].apply(lambda x: pd.to_numeric(x.split("%")[0]))

```

```

In [520... # checking the data types
loan.info()

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 39717 entries, 0 to 39716
Data columns (total 53 columns):
id                39717 non-null int64
member_id        39717 non-null int64
loan_amnt        39717 non-null int64
funded_amnt      39717 non-null int64
funded_amnt_inv  39717 non-null float64
term             39717 non-null object
int_rate         39717 non-null float64
installment      39717 non-null float64
grade            39717 non-null object
sub_grade        39717 non-null object
emp_title        37258 non-null object
emp_length       38642 non-null object
home_ownership   39717 non-null object
annual_inc       39717 non-null float64
verification_status 39717 non-null object
issue_d          39717 non-null object
loan_status      39717 non-null object
pymnt_plan       39717 non-null object
url              39717 non-null object
purpose          39717 non-null object
title            39706 non-null object
zip_code         39717 non-null object
addr_state       39717 non-null object
dti              39717 non-null float64
delinq_2yrs      39717 non-null int64
earliest_cr_line 39717 non-null object
inq_last_6mths   39717 non-null int64
open_acc         39717 non-null int64
pub_rec          39717 non-null int64
revol_bal        39717 non-null int64
revol_util       39667 non-null object
total_acc        39717 non-null int64
initial_list_status 39717 non-null object
out_prncp        39717 non-null float64
out_prncp_inv    39717 non-null float64
total_pymnt      39717 non-null float64
total_pymnt_inv  39717 non-null float64

```

```

total_rec_prncp      39717 non-null float64
total_rec_int        39717 non-null float64
total_rec_late_fee    39717 non-null float64
recoveries           39717 non-null float64
collection_recovery_fee 39717 non-null float64
last_pymnt_d         39646 non-null object
last_pymnt_amnt       39717 non-null float64
last_credit_pull_d    39715 non-null object
collections_12_mths_ex_med 39661 non-null float64
policy_code          39717 non-null int64
application_type      39717 non-null object
acc_now_delinq        39717 non-null int64
chargeoff_within_12_mths 39661 non-null float64
delinq_amnt          39717 non-null int64
pub_rec_bankruptcies  39020 non-null float64
tax_liens            39678 non-null float64
dtypes: float64(19), int64(13), object(21)
memory usage: 16.1+ MB

```

```

In [521... # also, lets extract the numeric part from the variable employment length

# first, let's drop the missing values from the column (otherwise the regex code below throws error)
loan = loan[~loan['emp_length'].isnull()]

# using regular expression to extract numeric values from the string
import re
loan['emp_length'] = loan['emp_length'].apply(lambda x: re.findall('\d+', str(x))[0])

# convert to numeric
loan['emp_length'] = loan['emp_length'].apply(lambda x: pd.to_numeric(x))

In [522... # Looking at type of the columns again
loan.info()

```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 38642 entries, 0 to 39716
Data columns (total 53 columns):
id                38642 non-null int64
member_id         38642 non-null int64
loan_amnt         38642 non-null int64
funded_amnt       38642 non-null int64
funded_amnt_inv   38642 non-null float64
term              38642 non-null object
int_rate          38642 non-null float64
installment       38642 non-null float64
grade             38642 non-null object
sub_grade         38642 non-null object
emp_title         37202 non-null object
emp_length        38642 non-null int64
home_ownership    38642 non-null object
annual_inc        38642 non-null float64
verification_status 38642 non-null object
issue_d           38642 non-null object
loan_status       38642 non-null object
pymnt_plan        38642 non-null object
url               38642 non-null object
purpose           38642 non-null object
title             38632 non-null object
zip_code          38642 non-null object
addr_state        38642 non-null object
dti               38642 non-null float64
delinq_2yrs       38642 non-null int64
earliest_cr_line  38642 non-null object
inq_last_6mths    38642 non-null int64
open_acc          38642 non-null int64
pub_rec           38642 non-null int64
revol_bal         38642 non-null int64
revol_util        38595 non-null object
total_acc         38642 non-null int64
initial_list_status 38642 non-null object
out_prncp         38642 non-null float64
out_prncp_inv     38642 non-null float64
total_pymnt       38642 non-null float64
total_pymnt_inv   38642 non-null float64

```



```

total_rec_prncp      38642 non-null float64
total_rec_int        38642 non-null float64
total_rec_late_fee   38642 non-null float64
recoveries           38642 non-null float64
collection_recovery_fee 38642 non-null float64
last_pymnt_d         38576 non-null object
last_pymnt_amnt      38642 non-null float64
last_credit_pull_d   38640 non-null object
collections_12_mths_ex_med 38586 non-null float64
policy_code          38642 non-null int64
application_type     38642 non-null object
acc_now_delinq       38642 non-null int64
chargeoff_within_12_mths 38586 non-null float64
delinq_amnt          38642 non-null int64
pub_rec_bankruptcies 37945 non-null float64
tax_liens            38603 non-null float64
dtypes: float64(19), int64(14), object(20)
memory usage: 15.9+ MB

```

Data Analysis

Let's now move to data analysis. To start with, let's understand the objective of the analysis clearly and identify the variables that we want to consider for analysis.

The objective is to identify predictors of default so that at the time of loan application, we can use those variables for approval/rejection of the loan. Now, there are broadly three types of variables - 1. those which are related to the applicant (demographic variables such as age, occupation, employment details etc.), 2. loan characteristics (amount of loan, interest rate, purpose of loan etc.) and 3. Customer behaviour variables (those which are generated after the loan is approved such as delinquent 2 years, revolving balance, next payment date etc.).

Now, the customer behaviour variables are not available at the time of loan application, and thus they cannot be used as predictors for credit approval.

Thus, going forward, we will use only the other two types of variables.

```

In [523... behaviour_var = [
    "delinq_2yrs",
    "earliest_cr_line",

```

```
"inq_last_6mths",  
"open_acc",  
"pub_rec",  
"revol_bal",  
"revol_util",  
"total_acc",  
"out_prncp",  
"out_prncp_inv",  
"total_pymnt",  
"total_pymnt_inv",  
"total_rec_prncp",  
"total_rec_int",  
"total_rec_late_fee",  
"recoveries",  
"collection_recovery_fee",  
"last_pymnt_d",  
"last_pymnt_amnt",  
"last_credit_pull_d",  
"application_type"]  
behaviour_var
```

```
Out[523]: ['delinq_2yrs',  
          'earliest_cr_line',  
          'inq_last_6mths',  
          'open_acc',  
          'pub_rec',  
          'revol_bal',  
          'revol_util',  
          'total_acc',  
          'out_prncp',  
          'out_prncp_inv',  
          'total_pymnt',  
          'total_pymnt_inv',  
          'total_rec_prncp',  
          'total_rec_int',  
          'total_rec_late_fee',  
          'recoveries',  
          'collection_recovery_fee',  
          'last_pymnt_d',  
          'last_pymnt_amnt',  
          'last_credit_pull_d',  
          'application_type']
```

```
In [524... # Let's now remove the behaviour variables from analysis  
df = loan.drop(behaviour_var, axis=1)  
df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 38642 entries, 0 to 39716
Data columns (total 32 columns):
id                38642 non-null int64
member_id         38642 non-null int64
loan_amnt         38642 non-null int64
funded_amnt       38642 non-null int64
funded_amnt_inv   38642 non-null float64
term              38642 non-null object
int_rate          38642 non-null float64
installment       38642 non-null float64
grade             38642 non-null object
sub_grade         38642 non-null object
emp_title         37202 non-null object
emp_length        38642 non-null int64
home_ownership    38642 non-null object
annual_inc        38642 non-null float64
verification_status 38642 non-null object
issue_d           38642 non-null object
loan_status       38642 non-null object
pymnt_plan        38642 non-null object
url               38642 non-null object
purpose           38642 non-null object
title             38632 non-null object
zip_code          38642 non-null object
addr_state        38642 non-null object
dti               38642 non-null float64
initial_list_status 38642 non-null object
collections_12_mths_ex_med 38586 non-null float64
policy_code       38642 non-null int64
acc_now_delinq    38642 non-null int64
chargeoff_within_12_mths 38586 non-null float64
delinq_amnt       38642 non-null int64
pub_rec_bankruptcies 37945 non-null float64
tax_liens         38603 non-null float64
dtypes: float64(9), int64(8), object(15)
memory usage: 9.7+ MB

```

Typically, variables such as acc_now_delinquent, chargeoff within 12 months etc. (which are related to the applicant's past loans) are available from the credit bureau.

```
In [525... # also, we will not be able to use the variables zip code, address, state etc.  
# the variable 'title' is derived from the variable 'purpose'  
# thus let get rid of all these variables as well  
  
df = df.drop(['title', 'url', 'zip_code', 'addr_state'], axis=1)
```

Next, let's have a look at the target variable - loan_status. We need to relabel the values to a binary form - 0 or 1, 1 indicating that the person has defaulted and 0 otherwise.

```
In [526... df['loan_status'] = df['loan_status'].astype('category')  
df['loan_status'].value_counts()
```

```
Out[526]: Fully Paid      32145  
Charged Off    5399  
Current        1098  
Name: loan_status, dtype: int64
```

You can see that fully paid comprises most of the loans. The ones marked 'current' are neither fully paid not defaulted, so let's get rid of the current loans. Also, let's tag the other two values as 0 or 1.

```
In [527... # filtering only fully paid or charged-off  
df = df[df['loan_status'] != 'Current']  
df['loan_status'] = df['loan_status'].apply(lambda x: 0 if x=='Fully Paid' else 1)  
  
# converting loan_status to integer type  
df['loan_status'] = df['loan_status'].apply(lambda x: pd.to_numeric(x))  
  
# summarising the values  
df['loan_status'].value_counts()
```

```
Out[527]: 0      32145  
1       5399  
Name: loan_status, dtype: int64
```

Next, let's start with univariate analysis and then move to bivariate analysis.

Univariate Analysis

First, let's look at the overall default rate.


```
In [528... # default rate
round(np.mean(df['loan_status']), 2)
```

```
Out[528]: 0.14000000000000001
```

The overall default rate is about 14%.


Let's first visualise the average default rates across categorical variables.

```
In [529... # plotting default rates across grade of the loan
sns.barplot(x='grade', y='loan_status', data=df)
plt.show()
```




```
In [1]: # Lets define a function to plot loan_status across categorical variables
def plot_cat(cat_var):
    sns.barplot(x=cat_var, y='loan_status', data=df, orient='v')
    plt.show()
```

```
In [531... # compare default rates across grade of loan  
plot_cat('grade')
```




Clearly, as the grade of loan goes from A to G, the default rate increases. This is expected because the grade is decided by Lending Club based on the riskiness of the loan.


```
In [532... # term: 60 months Loans default more than 36 months Loans  
plot_cat('term')
```




```
In [533... # sub-grade: as expected - A1 is better than A2 better than A3 and so on  
plt.figure(figsize=(16, 6))  
plot_cat('sub_grade')
```


```
In [534... # home ownership: not a great discriminator  
plot_cat('home_ownership')
```



```
In [535... # verification_status: surprisingly, verified loans default more than not verifiedb  
plot_cat('verification_status')
```



```
In [536... # purpose: small business loans default the most, then renewable energy and education  
plt.figure(figsize=(16, 6))  
plot_cat('purpose')
```



```
In [537... # Let's also observe the distribution of Loans across years
# first lets convert the year column into datetime and then extract year and month from it
df['issue_d'].head()
```

```
Out[537]: 0    Dec-11
1    Dec-11
2    Dec-11
3    Dec-11
5    Dec-11
Name: issue_d, dtype: object
```

```
In [538... from datetime import datetime
df['issue_d'] = df['issue_d'].apply(lambda x: datetime.strptime(x, '%b-%y'))
```

```
In [539... # extracting month and year from issue_date
df['month'] = df['issue_d'].apply(lambda x: x.month)
df['year'] = df['issue_d'].apply(lambda x: x.year)
```

```
In [540]: # Let's first observe the number of loans granted across years  
df.groupby('year').year.count()
```

```
Out[540]: year  
2007      251  
2008     1562  
2009     4716  
2010    11214  
2011    19801  
Name: year, dtype: int64
```


You can see that the number of loans has increased steadily across years.

```
In [541]: # number of Loans across months  
df.groupby('month').month.count()
```


```
Out[541]: month  
1      2331  
2      2278  
3      2632  
4      2756  
5      2838  
6      3094  
7      3253  
8      3321  
9      3394  
10     3637  
11     3890  
12     4120  
Name: month, dtype: int64
```

Most loans are granted in December, and in general in the latter half of the year.

```
In [542]: # Lets compare the default rates across years  
# the default rate had suddenly increased in 2011, inspite of reducing from 2008 till 2010  
plot_cat('year')
```




```
In [543... # comparing default rates across months: not much variation across months
plt.figure(figsize=(16, 6))
plot_cat('month')
```



Let's now analyse how the default rate varies across continuous variables.

```
In [544... # Loan amount: the median loan amount is around 10,000
sns.distplot(df['loan_amnt'])
plt.show()
```



The easiest way to analyse how default rates vary across continuous variables is to bin the variables into discrete categories.

Let's bin the loan amount variable into small, medium, high, very high.


```
In [545... # binning loan amount
def loan_amount(n):
    if n < 5000:
        return 'low'
    elif n >=5000 and n < 15000:
        return 'medium'
    elif n >= 15000 and n < 25000:
        return 'high'
    else:
        return 'very high'

df['loan_amnt'] = df['loan_amnt'].apply(lambda x: loan_amount(x))
```

```
In [546... df['loan_amnt'].value_counts()
```


```
Out[546]: medium      20157  
         high       7572  
         low        7095  
         very high   2720  
         Name: loan_amnt, dtype: int64
```

```
In [547... # Let's compare the default rates across loan amount type  
# higher the loan amount, higher the default rate  
plot_cat('loan_amnt')
```



```
In [548... # Let's also convert funded amount invested to bins  
df['funded_amnt_inv'] = df['funded_amnt_inv'].apply(lambda x: loan_amount(x))
```


```
In [549... # funded amount invested  
plot_cat('funded_amnt_inv')
```

```
In [550... # Lets also convert interest rate to low, medium, high
# binning loan amount
def int_rate(n):
    if n <= 10:
        return 'low'
    elif n > 10 and n <=15:
        return 'medium'
    else:
        return 'high'

df['int_rate'] = df['int_rate'].apply(lambda x: int_rate(x))
```


```
In [551... # comparing default rates across rates of interest
# high interest rates default more, as expected
plot_cat('int_rate')
```



```
In [552... # debt to income ratio
def dti(n):
    if n <= 10:
        return 'low'
    elif n > 10 and n <=20:
        return 'medium'
    else:
        return 'high'

df['dti'] = df['dti'].apply(lambda x: dti(x))
```


```
In [553... # comparing default rates across debt to income ratio
# high dti translates into higher default rates, as expected
plot_cat('dti')
```



```
In [554... # Segregating the funded amount into the respective bins into Low, Medium and High
def funded_amount(n):
    if n <= 5000:
        return 'low'
    elif n > 5000 and n <=15000:
        return 'medium'
    else:
        return 'high'

df['funded_amnt'] = df['funded_amnt'].apply(lambda x: funded_amount(x))
```


```
In [555... plot_cat('funded_amnt')
```



```
In [556... # installment dividing into the respective bins into Low, Medium, High and very High
def installment(n):
    if n <= 200:
        return 'low'
    elif n > 200 and n <=400:
        return 'medium'
    elif n > 400 and n <=600:
        return 'high'
    else:
        return 'very high'

df['installment'] = df['installment'].apply(lambda x: installment(x))
```


```
In [557... # comparing default rates across installment
# the higher the installment amount, the higher the default rate
plot_cat('installment')
```



```
In [558... # annual income dividing into the respective bins into Low, Medium, High and very High
def annual_income(n):
    if n <= 50000:
        return 'low'
    elif n > 50000 and n <=100000:
        return 'medium'
    elif n > 100000 and n <=150000:
        return 'high'
    else:
        return 'very high'

df['annual_inc'] = df['annual_inc'].apply(lambda x: annual_income(x))
```

```
In [559... # annual income and default rate
# Lower the annual income, higher the default rate
plot_cat('annual_inc')
```




```
In [560... # employment length
# first, let's drop the missing value observations in emp length
df = df[~df['emp_length'].isnull()]

# binning the variable
def emp_length(n):
    if n <= 1:
        return 'fresher'
    elif n > 1 and n <=3:
        return 'junior'
    elif n > 3 and n <=7:
        return 'senior'
    else:
        return 'expert'

df['emp_length'] = df['emp_length'].apply(lambda x: emp_length(x))
```

```
In [561... # emp_length and default rate
# not much of a predictor of default
plot_cat('emp_length')
```




Segmented Univariate Analysis

We have now compared the default rates across various variables, and some of the important predictors are purpose of the loan, interest rate, annual income, grade etc.

In the credit industry, one of the most important factors affecting default is the purpose of the loan - home loans perform differently than credit cards, credit cards are very different from debt consolidation loans etc.


This comes from business understanding, though let's again have a look at the default rates across the purpose of the loan.

```
In [562... # purpose: small business loans default the most, then renewable energy and education
plt.figure(figsize=(16, 6))
plot_cat('purpose')
```



In the upcoming analyses, we will segment the loan applications across the purpose of the loan, since that is a variable affecting many other variables - the type of applicant, interest rate, income, and finally the default rate.

```
In [563... # Lets first look at the number of loans for each type (purpose) of the loan
# most loans are debt consolidation (to repay otehr debts), then credit card, major purchase etc.
plt.figure(figsize=(16, 6))
sns.countplot(x='purpose', data=df)
plt.show()
```





Let's analyse the top 4 types of loans based on purpose: consolidation, credit card, home improvement and major purchase.

```
In [564... # filtering the df for the 4 types of loans mentioned above
main_purposes = ["credit_card", "debt_consolidation", "home_improvement", "major_purchase"]
df = df[df['purpose'].isin(main_purposes)]
df['purpose'].value_counts()
```


```
Out[564]: debt_consolidation    17675
credit_card                    4899
home_improvement              2785
major_purchase                2080
Name: purpose, dtype: int64
```

```
In [565... # plotting number of Loans by purpose
sns.countplot(x=df['purpose'])
plt.show()
```



In [566... *# Let's now compare the default rates across two types of categorical variables*
purpose of loan (constant) and another categorical variable (which changes)


```
plt.figure(figsize=[10, 6])  
sns.barplot(x='term', y="loan_status", hue='purpose', data=df)  
plt.show()
```




In [567... *# Lets write a function which takes a categorical variable and plots the default rate
segmented by purpose*

```
def plot_segmented(cat_var):  
    plt.figure(figsize=(10, 6))  
    sns.barplot(x=cat_var, y='loan_status', hue='purpose', data=df)  
    plt.show()
```


```
plot_segmented('term')
```



```
In [568... # grade of Loan  
plot_segmented('grade')
```




```
In [569... # home ownership  
plot_segmented('home_ownership')
```




In general, debt consolidation loans have the highest default rates. Lets compare across other categories as well.

In [570...


```
# year  
plot_segmented('year')
```




```
In [571... # emp_length  
plot_segmented('emp_length')
```




```
In [572... # loan_amnt: same trend across loan purposes  
plot_segmented('loan_amnt')
```


```
In [573... # interest rate  
plot_segmented('int_rate')
```




```
In [574... # installment  
plot_segmented('installment')
```



```
In [575... # debt to income ratio  
plot_segmented('dti')
```



```
In [576... # annual income  
plot_segmented('annual_inc')
```



A good way to quantify the effect of a categorical variable on default rate is to see 'how much does the default rate vary across the categories'.

Let's see an example using annual_inc as the categorical variable.

```
In [577]: # variation of default rate across annual_inc
df.groupby('annual_inc').loan_status.mean().sort_values(ascending=False)
```

```
Out[577]: annual_inc
low      0.157966
medium   0.130075
very high 0.101570
high      0.097749
Name: loan_status, dtype: float64
```

```
In [578]: # one can write a function which takes in a categorical variable and computed the average
# default rate across the categories
# It can also compute the 'difference between the highest and the lowest default rate' across the
```

```
# categories, which is a decent metric indicating the effect of the variable on default rate

def diff_rate(cat_var):
    default_rates = df.groupby(cat_var).loan_status.mean().sort_values(ascending=False)
    return (round(default_rates, 2), round(default_rates[0] - default_rates[-1], 2))

default_rates, diff = diff_rate('annual_inc')
print(default_rates)
print(diff)
```

```
annual_inc
low      0.16
medium   0.13
very high 0.10
high     0.10
Name: loan_status, dtype: float64
0.06
```

Thus, there is a 6% increase in default rate as you go from high to low annual income. We can compute this difference for all the variables and roughly identify the ones that affect default rate the most.

In [579...

```
# filtering all the object type variables
df_categorical = df.loc[:, df.dtypes == object]
df_categorical['loan_status'] = df['loan_status']

# Now, for each variable, we can compute the incremental diff in default rates
print([i for i in df.columns])
```

```
['id', 'member_id', 'loan_amnt', 'funded_amnt', 'funded_amnt_inv', 'term', 'int_rate', 'installment', 'grade', 'sub_grade', 'emp_title', 'emp_length', 'home_ownership', 'annual_inc', 'verification_status', 'issue_d', 'loan_status', 'pymnt_plan', 'purpose', 'dti', 'initial_list_status', 'collections_12_mths_ex_med', 'policy_code', 'acc_now_delinq', 'chargeoff_within_12_mths', 'delinq_amnt', 'pub_rec_bankruptcies', 'tax_liens', 'month', 'year']
```

```
/Library/Frameworks/Python.framework/Versions/3.5/lib/python3.5/site-packages/ipykernel_launcher.py:3: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy>
This is separate from the ipykernel package so we can avoid doing imports until

In [596...

storing the diff of default rates for each column in a dict

```
d = {key: diff_rate(key)[1]*100 for key in df_categorical.columns if key != 'loan_status'}  
print(d)
```

```
{'loan_amnt': 7.0000000000000009, 'funded_amnt_inv': 6.0, 'pymnt_plan': 0.0, 'verification_status': 4.0, 'emp_title': 100.0, 'dti': 5.0, 'home_ownership': 16.0, 'purpose': 5.0, 'sub_grade': 46.0, 'grade': 27.0, 'funded_amnt': 5.0, 'installment': 3.0, 'initial_list_status': 0.0, 'int_rate': 19.0, 'term': 15.0, 'annual_inc': 6.0, 'emp_length': 2.0}
```